

Story 1.3: Kafka Infrastructure & Schema Registry

Status: ready-for-dev

Story

As a Developer,

I want Kafka cluster con Schema Registry configurado para eventos Avro,
so that Puedo publicar domain events con garantía de schema versionado y backward compatibility.

Acceptance Criteria

AC1: Kafka + Zookeeper + Schema Registry Docker Compose

Given El proyecto tiene Docker Compose configurado

When Agrego servicios de Kafka al `docker-compose.yml`

Then

- Servicio `zookeeper` configurado:
 - Imagen: `confluentinc/cp-zookeeper:7.5.0`
 - Puerto: `2181:2181`
 - Configuración: `ZOOKEEPER_CLIENT_PORT=2181`
- Servicio `kafka` configurado:
 - Imagen: `confluentinc/cp-kafka:7.5.0`
 - Puerto: `9092:9092 (external)`, `29092:29092 (internal)`
 - Bootstrap servers: `localhost:9092 (dev)`, `kafka:29092 (container-to-container)`
 - Depende de `zookeeper`
- Servicio `schema-registry` configurado:
 - Imagen: `confluentinc/cp-schema-registry:7.5.0`
 - Puerto: `8081:8081`
 - URL: `http://localhost:8081`
 - Depende de `kafka`
- Comando `docker-compose up -d` levanta los 3 servicios exitosamente
- Healthchecks configurados para Kafka y Schema Registry

AC2: Spring Kafka Dependencies

Given El proyecto tiene Spring Boot 3.2+

When Agrego dependencias de Kafka

Then

- pom.xml incluye:
 - spring-kafka (versión gestionada por Spring Boot)
 - kafka-avro-serializer (Confluent 7.5.0)
 - avro (Apache Avro 1.11+)
 - kafka-streams-test-utils (test scope)
 - spring-kafka-test (test scope)

AC3: Kafka Configuration (Spring Boot)

Given Kafka running en Docker

When Configuro application-local.yml

Then

- Producer configuration:
 - spring.kafka.bootstrap-servers=localhost:9092
 - spring.kafka.producer.key-
serializer=org.apache.kafka.common.serialization.StringSerializer
 - spring.kafka.producer.value-
serializer=io.confluent.kafka.serializers.KafkaAvroSerializer
 - spring.kafka.producer.acks=all (strong durability)
 - spring.kafka.producer.compression-type=snappy
 - spring.kafka.producer.max-in-flight-requests-per-connection=5
- Schema Registry configuration:
 - spring.kafka.properties.schema.registry.url=http://localhost:8081
- Admin configuration:
 - spring.kafka.admin.auto-create=true

AC4: Avro Schema Definition

Given Schema Registry configurado

When Defino esquema Avro para eventos de dominio

Then

- Archivo `src/main/resources/kafka/schemas/signature-event.avsc` creado
- Esquema define:
 - namespace: com.bank.signature.event
 - name: SignatureEvent
 - type: record
 - Campos comunes: eventId, eventType, aggregateId, aggregateType, timestamp, traceId
 - Campo payload: Union type con 8 event types:
 - SIGNATURE_REQUEST_CREATED
 - CHALLENGE_SENT
 - CHALLENGE_COMPLETED
 - CHALLENGE_FAILED
 - SIGNATURE_COMPLETED
 - SIGNATURE_FAILED
 - FALLBACK_TRIGGERED
 - PROVIDER_DEGRADED
 - Backward compatibility validada

AC5: Kafka Topic Creation

Given Kafka Admin configurado con `auto-create=true`

When La aplicación inicia

Then

- Topic `signature.events` creado automáticamente:
 - Partitions: 12 (para throughput)
 - Replication factor: 1 (dev), 3 (prod)
 - Retention: 7 días
 - Compression: snappy
- Topic `signature.events.dlq` (Dead Letter Queue) creado:
 - Partitions: 3
 - Replication factor: 1 (dev), 3 (prod)
 - Retention: 30 días

AC6: KafkaTemplate Configuration

Given Spring Kafka configurado

When Creo KafkaConfig.java

Then

- Bean KafkaTemplate<String, GenericRecord> configurado
- ProducerFactory con:
 - Key serializer: StringSerializer
 - Value serializer: KafkaAvroSerializer
 - Idempotence habilitado (enable.idempotence=true)
 - Transactional ID configurado (para exactly-once semantics en futuro)
- Default topic: signature.events

AC7: Schema Registration in Schema Registry

Given Schema Registry running

When Registro esquema Avro

Then

- Esquema signature-event-value registrado en Schema Registry
- Subject: signature.events-value (key strategy: TopicNameStrategy)
- Compatibility mode: BACKWARD (permite agregar campos opcionales)
- Schema ID asignado (e.g., 1)
- GET `http://localhost:8081/subjects` retorna ["signature.events-value"]
- GET `http://localhost:8081/subjects/signature.events-value/versions/latest` retorna schema

AC8: Kafka Health Check

Given Kafka configurado en Spring Boot

When Configuro Actuator health check

Then

- Endpoint /actuator/health/kafka retorna {"status": "UP", "details": {"kafkaConsumers": "UP", "kafkaProducers": "UP"} }
- Health check verifica:
 - Conexión a Kafka broker exitosa
 - Producer está listo para enviar mensajes

- Si Kafka está down, endpoint retorna { "status": "DOWN" }

AC9: Maven Avro Plugin Configuration

Given Esquema Avro definido en `.avsc`

When Configuro `avro-maven-plugin` en `pom.xml`

Then

- Plugin `avro-maven-plugin` configurado en `<build><plugins>`
- Goal: `schema` (genera clases Java desde `.avsc`)
- Source directory: `src/main/resources/kafka/schemas`
- Output directory: `target/generated-sources/avro`
- Comando `./mvnw clean compile` genera clases:
 - `com.bank.signature.event.SignatureEvent`
 - Builders, getters, setters para cada event type
- Clases generadas disponibles en classpath

AC10: Integration Test with Embedded Kafka

Given Spring Kafka Test configurado

When Creo test de integración `KafkaInfrastructureIntegrationTest.java`

Then

- Test usa `@EmbeddedKafka` con:
 - Topics: `signature.events`, `signature.events.dlq`
 - Partitions: 3 (embedded)
 - Broker properties: auto-create topics, port 9093
- Test verifica:
 - `KafkaTemplate` puede enviar mensaje `GenericRecord`
 - Mensaje se serializa correctamente con Avro
 - Schema Registry (mock) valida esquema
 - Mensaje llega a topic `signature.events`
- Test pasa en `mvn verify`

AC11: Kafka Configuration Profiles

Given Múltiples entornos (local, uat, prod)

When Configuro profiles en `application-{profile}.yml`

Then

- `application-local.yml`:
 - `bootstrap-servers: localhost:9092`
 - `schema.registry.url: http://localhost:8081`
- `application-uat.yml (futuro)`:
 - `bootstrap-servers: kafka-uat.internal:9092`
 - `schema.registry.url: http://schema-registry-uat.internal:8081`
- `application-prod.yml (futuro)`:
 - `bootstrap-servers: kafka-prod.internal:9092`
 - `schema.registry.url: http://schema-registry-prod.internal:8081`
 - `producer.acks: all`
 - `producer.enable.idempotence: true`

AC12: Documentation & README Update

Given Kafka infrastructure configurado

When Actualizo documentación

Then

- `README.md` actualizado con sección "Kafka Setup":
 - Comandos Docker Compose para Kafka
 - Comandos para verificar topics: `docker exec kafka kafka-topics --bootstrap-server localhost:9092 --list`
 - Comandos para Schema Registry: `curl http://localhost:8081/subjects`
- `docs/development/kafka-messaging.md` creado con:
 - Avro schema evolution guidelines
 - Event publishing patterns
 - Testing strategy con Embedded Kafka
 - Troubleshooting (Kafka connection errors, schema validation failures)
- `CHANGELOG.md` actualizado con Story 1.3 entry

Tasks / Subtasks

Task 1: Add Kafka Services to Docker Compose (AC: #1)

- 1.1. Agregar servicio `zookeeper` a `docker-compose.yml`:
 - Imagen: `confluentinc/cp-zookeeper:7.5.0`
 - Puerto: `2181`
 - Variable: `ZOOKEEPER_CLIENT_PORT=2181`
- 1.2. Agregar servicio `kafka` a `docker-compose.yml`:
 - Imagen: `confluentinc/cp-kafka:7.5.0`
 - Puertos: `9092` (external), `29092` (internal)
 - Variables: `KAFKA_BROKER_ID`, `KAFKA_ZOOKEEPER_CONNECT`,
`KAFKA_ADVERTISED_LISTENERS`, `KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR`
 - Depende de: `zookeeper`
- 1.3. Agregar servicio `schema-registry` a `docker-compose.yml`:
 - Imagen: `confluentinc/cp-schema-registry:7.5.0`
 - Puerto: `8081`
 - Variables: `SCHEMA_REGISTRY_HOST_NAME`,
`SCHEMA_REGISTRY_KAFKASTORE_BOOTSTRAP_SERVERS`
 - Depende de: `kafka`
- 1.4. Agregar healthchecks para Kafka y Schema Registry
- 1.5. Verificar: `docker-compose up -d` levanta los 3 servicios

Task 2: Add Spring Kafka Dependencies to pom.xml (AC: #2)

- 2.1. Agregar `spring-kafka` (Spring Boot managed version)
- 2.2. Agregar `kafka-avro-serializer` (Confluent 7.5.0)
- 2.3. Agregar `avro` (Apache Avro 1.11+)
- 2.4. Agregar `kafka-streams-test-utils` (test scope)
- 2.5. Agregar `spring-kafka-test` (test scope)

Task 3: Configure Kafka in application-local.yml (AC: #3)

- 3.1. Configurar `spring.kafka.bootstrap-servers=localhost:9092`
- 3.2. Configurar producer serializers (`String + KafkaAvroSerializer`)
- 3.3. Configurar producer properties: `acks=all, compression-type=snappy, max-in-`

flight=5

- 3.4. Configurar Schema Registry URL: `http://localhost:8081`
- 3.5. Configurar admin auto-create: `spring.kafka.admin.auto-create=true`

Task 4: Define Avro Schema for Domain Events (AC: #4)

- 4.1. Crear directorio `src/main/resources/kafka/schemas/`
- 4.2. Crear archivo `signature-event.avsc` con namespace `com.bank.signature.event`
- 4.3. Definir campos comunes: `eventId, eventType, aggregateId, timestamp, traceId`
- 4.4. Definir union type `payload` con 8 event types (SIGNATURE_REQUEST_CREATED, CHALLENGE_SENT, etc.)
- 4.5. Validar esquema con Avro tools: `java -jar avro-tools.jar compile schema signature-event.avsc .`

Task 5: Configure Kafka Topics (AC: #5)

- 5.1. Crear `KafkaTopicConfig.java` con beans `NewTopic`
- 5.2. Configurar topic `signature.events`:
 - 12 partitions, replication factor 1 (dev)
 - Retention 7 días, compression snappy
- 5.3. Configurar topic `signature.events.dlq`:
 - 3 partitions, replication factor 1 (dev)
 - Retention 30 días
- 5.4. Verificar topics con: `docker exec kafka kafka-topics --bootstrap-server localhost:9092 --list`

Task 6: Configure KafkaTemplate Bean (AC: #6)

- 6.1. Crear `KafkaConfig.java` en `infrastructure/config/`
- 6.2. Configurar `ProducerFactory<String, GenericRecord>` con `KafkaAvroSerializer`
- 6.3. Configurar `KafkaTemplate<String, GenericRecord>` bean
- 6.4. Habilitar idempotence: `enable.idempotence=true`
- 6.5. Configurar default topic: `signature.events`

Task 7: Register Schema in Schema Registry (AC: #7)

- 7.1. Iniciar Schema Registry: `docker-compose up -d schema-registry`
- 7.2. Configurar compatibility mode: BACKWARD
- 7.3. Registrar schema vía API o auto-registro en primer envío
- 7.4. Verificar: `curl http://localhost:8081/subjects`
- 7.5. Verificar schema: `curl http://localhost:8081/subjects/signature.events-value/versions/latest`

Task 8: Configure Kafka Health Check (AC: #8)

- 8.1. Verificar `spring-boot-starter-actuator` está en `pom.xml`
- 8.2. Configurar `management.health.kafka.enabled=true` en `application.yml`
- 8.3. Exponer endpoint en `management.endpoints.web.exposure.include` (ya configurado)
- 8.4. Verificar: `curl http://localhost:8080/actuator/health/kafka`
- 8.5. Test: detener Kafka, verificar health check retorna DOWN

Task 9: Configure Maven Avro Plugin (AC: #9)

- 9.1. Agregar `avro-maven-plugin` a `pom.xml` en `<build><plugins>`
- 9.2. Configurar goal `schema` para generar clases Java
- 9.3. Configurar source directory: `src/main/resources/kafka/schemas`
- 9.4. Configurar output directory: `target/generated-sources/avro`
- 9.5. Ejecutar: `./mvnw clean compile`
- 9.6. Verificar clases generadas: `com.bank.signature.event.SignatureEvent`

Task 10: Create Integration Test with Embedded Kafka (AC: #10)

- 10.1. Crear `KafkaInfrastructureIntegrationTest.java` en
`src/test/java/.../infrastructure/`
- 10.2. Configurar `@EmbeddedKafka` con topics: `signature.events`,
`signature.events.dlq`
- 10.3. Autowire `KafkaTemplate<String, GenericRecord>`
- 10.4. Test method: `testKafkaTemplateSendsAvroMessage()`
 - Crear `GenericRecord` con evento `SIGNATURE_REQUEST_CREATED`
 - Enviar con `KafkaTemplate`

- Consumir mensaje con KafkaConsumer
- Verificar serialización Avro correcta

10.5. Ejecutar: `./mvnw verify`

Task 11: Configure Kafka Profiles for Multiple Environments (AC: #11)

- 11.1. Configurar `application-local.yml` con bootstrap-servers localhost
- 11.2. Crear `application-uat.yml` con Kafka internal URLs (placeholder)
- 11.3. Crear `application-prod.yml` con Kafka internal URLs + idempotence (placeholder)
- 11.4. Documentar diferencias en `docs/development/kafka-messaging.md`

Task 12: Update Documentation (AC: #12)

- 12.1. Actualizar `README.md` con sección "Kafka Setup":
 - Comandos Docker Compose
 - Comandos para listar topics
 - Comandos Schema Registry (`curl`)
- 12.2. Crear `docs/development/kafka-messaging.md`:
 - Avro schema evolution guidelines
 - Event publishing patterns
 - Testing strategy
 - Troubleshooting
- 12.3. Actualizar `CHANGELOG.md` con Story 1.3 entry
- 12.4. Agregar comentarios en `KafkaConfig.java` explicando configuraciones

Dev Notes

Architecture Patterns & Constraints

- **Event-Driven Architecture:** Kafka como backbone para eventos de dominio (Outbox pattern en Story 1.2)
- **Avro Serialization:** Esquema versionado con backward compatibility garantiza evolución segura
- **Schema Registry:** Confluent Schema Registry valida esquemas antes de publicar (fail fast)
- **Idempotent Producer:** `enable.idempotence=true` garantiza exactly-once delivery en caso de retries
- **Partitioning Strategy:** 12 partitions para `signature.events` permite throughput alto (parallel consumers)

Source Tree Components to Touch

```
signature-router/
├── pom.xml                                     # [MODIFY] agregar spring-kafka,
kafka-avro-serializer, avro, avro-maven-plugin
├── docker-compose.yml                           # [MODIFY] agregar zookeeper,
kafka, schema-registry services
├── src/main/resources/
|   ├── application.yml                         # [MODIFY] kafka health check
config
|   ├── application-local.yml                  # [MODIFY] kafka bootstrap-
servers, schema registry URL
|   └── kafka/schemas/
|       └── signature-event.avsc            # [CREATE] Avro schema definition
└── src/main/java/com/bank/signature/infrastructure/config/
    └── KafkaConfig.java                     # [CREATE] KafkaTemplate,
ProducerFactory beans
    └── KafkaTopicConfig.java               # [CREATE] NewTopic beans
(signature.events, dlq)
└── src/test/java/com/bank/signature/infrastructure/
    └── KafkaInfrastructureIntegrationTest.java # [CREATE] @EmbeddedKafka test
└── target/generated-sources/avro/              # [AUTO-GENERATED] Avro classes
(mvn compile)
    └── com/bank/signature/event/
        └── SignatureEvent.java
└── docs/development/
    └── kafka-messaging.md                  # [CREATE] Kafka documentation
└── README.md                                  # [MODIFY] Kafka setup section
```

Testing Standards Summary

- **Unit Tests:** No aplicable (Kafka configuration es infrastructure setup)

- **Integration Tests:**

- KafkaInfrastructureIntegrationTest.java con @EmbeddedKafka
- Verifica KafkaTemplate puede enviar GenericRecord
- Verifica serialización Avro correcta
- Verifica Schema Registry (mock) valida esquema

- **Manual Tests:**

- docker-compose up -d levanta Kafka cluster
- curl http://localhost:8081/subjects lista schemas
- /actuator/health/kafka retorna UP

- **CI/CD Pipeline:**

- Docker Compose up en pipeline
- mvn verify (incluye Embedded Kafka tests)
- Validar health check UP

Project Structure Notes

- **Avro Schema Evolution:** BACKWARD compatibility permite agregar campos opcionales sin romper consumers
- **Event Types:** 8 event types definidos (alineados con `docs/architecture/04-event-catalog.md`)
- **Partitioning:** `aggregateId` (`signature_request.id`) como partition key garantiza orden por request
- **DLQ (Dead Letter Queue):** `signature.events.dlq` para mensajes fallidos (retry exhausted)
- **Schema Registry Subject:** `TopicNameStrategy → signature.events-value` (un schema por topic)

References

- [Source: [docs/architecture/04-event-catalog.md](#)]: Catálogo de 8 eventos de dominio
 - SIGNATURE_REQUEST_CREATED, CHALLENGE_SENT, CHALLENGE_COMPLETED, CHALLENGE_FAILED, SIGNATURE_COMPLETED, SIGNATURE_FAILED, FALLBACK_TRIGGERED, PROVIDER_DEGRADED
- [Source: [docs/sprint-artifacts/tech-spec-epic-1.md](#)]: Kafka technology stack
 - Kafka 3.6 (Confluent), Schema Registry 7.5, Avro serialization
- [Source: [docs/epics.md](#)]: Story 1.3 acceptance criteria
 - Topics: `signature.events` (12 partitions), `signature.events.dlq`
 - Producer: `acks=all`, `compression=snappy`
- [Source: [docs/prd.md](#)]: Event Publishing requirements (FR39-FR46)
 - Atomicidad (Outbox pattern - Story 1.2), serialización Avro, partitioning por `aggregate_id`

Critical Implementation Notes

- **Kafka Advertised Listeners:** Docker Compose debe configurar `KAFKA_ADVERTISED_LISTENERS` con `PLAINTEXT://localhost:9092, PLAINTEXT_INTERNAL://kafka:29092` para que app en host y containers puedan conectarse
- **Schema Registry Compatibility:** BACKWARD mode permite agregar campos opcionales,

eliminar campos con defaults

- **Idempotence:** `enable.idempotence=true + acks=all` garantiza exactly-once semantics (no duplicados en caso de retry)
- **Avro Maven Plugin:** Genera clases Java en `target/generated-sources/avro/`, debe agregarse a classpath (maven-compiler-plugin source path)
- **Embedded Kafka Test:** Puerto debe ser diferente (9093) para no colisionar con Kafka en Docker (9092)

Definition of Done

Code Complete:

- 3 servicios agregados a `docker-compose.yml` (zookeeper, kafka, schema-registry)
- 5 dependencies agregadas a `pom.xml` (spring-kafka, kafka-avro-serializer, avro, test utils)
- Avro schema `signature-event.avsc` definido con 8 event types
- `KafkaConfig.java` y `KafkaTopicConfig.java` creados
- `KafkaInfrastructureIntegrationTest.java` creado con `@EmbeddedKafka`
- Maven Avro Plugin configurado en `pom.xml`
- `application-local.yml` configurado con Kafka properties

Tests Passing:

- Integration test con `@EmbeddedKafka` pasa en `mvn verify`
- Manual test: `docker-compose up -d` levanta Kafka cluster exitosamente
- Manual test: `curl http://localhost:8081/subjects` lista schemas
- Manual test: `/actuator/health/kafka` retorna UP

Architecture Validated:

- Avro schema sigue naming conventions (namespace: `com.bank.signature.event`)
- Topics configurados con partitioning strategy (12 partitions para throughput)
- Idempotent producer habilitado (`enable.idempotence=true`)
- DLQ topic configurado para mensajes fallidos

Documentation Updated:

- `README.md` incluye sección "Kafka Setup" con comandos Docker Compose, verificación topics/schemas
- `docs/development/kafka-messaging.md` creado con guidelines de schema evolution, testing, troubleshooting

- KafkaConfig.java tiene comentarios explicando configuraciones (acks, idempotence, compression)
- CHANGELOG.md actualizado: "Added Kafka 3.6 + Schema Registry 7.5 with Avro serialization"

Code Review Approved:

- Peer review confirma Avro schema es backward compatible
- Verificar Kafka advertised listeners configurados correctamente (localhost + internal)
- Validar idempotence + acks=all para exactly-once semantics
- Confirmar topics tienen retention policies correctas (7 días events, 30 días DLQ)

Story Marked as Done:

- Todos los 12 Acceptance Criteria verificados
 - Sprint status actualizado: 1-3-kafka-infrastructure-schema-registry: done
 - Story list actualizada en `docs/sprint-artifacts/sprint-status.yaml`
-

Dev Agent Record

Context Reference

- `docs/sprint-artifacts/1-3-kafka-infrastructure-schema-registry.context.xml`

Agent Model Used

Claude Sonnet 4.5

Debug Log References

Completion Notes List

File List

Created:

Modified:

Deleted:

Change Log

Date	Author	Change
2025-11-26	BMAD SM Agent	Story 1.3 draft created with Kafka + Schema Registry + Avro
2025-11-26	BMAD SM Agent	Technical context generated, status: ready-for-dev